

```
In [1]: import pandas as pd
import numpy as np
%matplotlib inline
```

```
In [2]: df = pd.read_excel('C://Users//saswa//OneDrive//Desktop//Pinaki-Time-series-forecasting//Superstore_Sales_Records.xls', index_
df = df[df['Category']=='Furniture']
df = df.groupby(by='Order Date').agg({'Sales':sum})
df.sort_index(inplace=True)
df.head(15)
```

Out[2]:

Sales**Order Date**

2014-01-06	2573.820
2014-01-07	76.728
2014-01-10	51.940
2014-01-11	9.940
2014-01-13	879.939
2014-01-14	61.960
2014-01-16	127.104
2014-01-19	181.470
2014-01-20	1413.510
2014-01-21	25.248
2014-01-26	217.200
2014-01-27	333.000
2014-01-31	290.666
2014-02-08	14.560
2014-02-11	1650.050

In [3]: `df = df.resample('MS').sum()``df.head(15)`

Out[3]:

Sales	
Order Date	
2014-01-01	6242.5250
2014-02-01	1839.6580
2014-03-01	14573.9560
2014-04-01	7944.8370
2014-05-01	6912.7870
2014-06-01	13206.1256
2014-07-01	10821.0510
2014-08-01	7320.3465
2014-09-01	23816.4808
2014-10-01	12304.2470
2014-11-01	21564.8727
2014-12-01	30645.9665
2015-01-01	11739.9416
2015-02-01	3134.3740
2015-03-01	12499.7830

In [4]: `df.shape`Out[4]: `(48, 1)`

```
In [5]: n = len(df)
m = int(n*0.8)

train_data = df.iloc[0:m]
```

```
test_data = df.iloc[m:n]

print(f"Total df size {len(df)}")
print(f"Total train data size {len(train_data)}")
print(f"Total test data size {len(test_data)}")
```

Total df size 48
Total train data size 38
Total test data size 10

In [6]: `train_data.tail()`

Out[6]:

	Sales
Order Date	
2016-10-01	11872.5770
2016-11-01	31783.6288
2016-12-01	36678.7150
2017-01-01	5964.0320
2017-02-01	6866.3374

In [7]: `from statsmodels.tsa.holtwinters import ExponentialSmoothing`


```
fitted_model = ExponentialSmoothing(train_data['Sales'], trend='mul', seasonal='mul', seasonal_periods=12).fit()
```

C:\Users\saswa\AppData\Roaming\Python\Python311\site-packages\statsmodels\tsa\holtwinters\model.py:83: RuntimeWarning: overflow encountered in matmul
return err.T @ err
C:\Users\saswa\AppData\Roaming\Python\Python311\site-packages\statsmodels\tsa\holtwinters\model.py:917: ConvergenceWarning: Optimization failed to converge. Check mle_retvals.
warnings.warn(

In [8]: `test_predictions = fitted_model.forecast(len(test_data))`
`test_predictions`


```
Out[8]: 2017-03-01    2.139283e+18
        2017-04-01    4.589978e+18
        2017-05-01    1.088440e+19
        2017-06-01    2.139505e+19
        2017-07-01    6.488781e+19
        2017-08-01    1.055809e+20
        2017-09-01    7.192510e+20
        2017-10-01    8.114941e+20
        2017-11-01    4.000376e+21
        2017-12-01    9.608036e+21
        Freq: MS, dtype: float64
```

```
In [9]: train_data['Sales'].plot(legend=True, label='TRAIN DATA')
        test_data['Sales'].plot(legend=True, label='TEST DATA', figsize=(12, 6)).autoscale(axis='x', tight=True)
```



```
In [10]: train_data['Sales'].plot(legend=True, label='TRAIN DATA')
test_data['Sales'].plot(legend=True, label='TEST DATA', figsize=(12, 6)).autoscale(axis='x', tight=True)
test_predictions.plot(legend=True, label='PREDICTION')
```

```
Out[10]: <Axes: xlabel='Order Date'>
```



```
In [11]: from sklearn.metrics import mean_absolute_error, mean_squared_error
mae_error = mean_absolute_error(test_data, test_predictions)

print(f"Mean absolute error of the above model is {mae_error}")
```

Mean absolute error of the above model is $1.5348634368862705 \times 10^{21}$

```
In [12]: mse_error = mean_squared_error(test_data, test_predictions)
```

```
print(f"Mean squared error of the above model is {mse_error}")
```

Mean squared error of the above model is 1.0950916485965921e+43

```
In [13]: rmse_error = np.sqrt(mean_squared_error(test_data, test_predictions))
```

```
print(f"Root mean squared error of the above model is {rmse_error}")
```

Root mean squared error of the above model is 3.309216899202275e+21

```
In [14]: test_data.describe()
```

```
Out[14]:
```

	Sales
count	10.000000
mean	20255.689980
std	9463.329001
min	9065.958100
25%	12720.235000
50%	17983.072450
75%	27242.171550
max	37056.715000

```
In [15]: fitted_model = ExponentialSmoothing(df['Sales'], trend='mul', seasonal='mul', seasonal_periods=12).fit()
```

```
C:\Users\saswa\AppData\Roaming\Python\Python311\site-packages\statsmodels\tsa\holtwinters\model.py:917: ConvergenceWarning: Optimization failed to converge. Check mle_retvals.
warnings.warn(
```


```
In [16]: future_preds = fitted_model.forecast(12)
future_preds
```



```
Out[16]: 2018-01-01    10145.196758
          2018-02-01     5276.646045
          2018-03-01    14046.143518
          2018-04-01    12638.938600
          2018-05-01    15347.806822
          2018-06-01    14398.498944
          2018-07-01    15575.177149
          2018-08-01    12050.471323
          2018-09-01    32198.645769
          2018-10-01    15228.732215
          2018-11-01    34823.955051
          2018-12-01    37118.533234
          Freq: MS, dtype: float64
```

```
In [17]: df['Sales'].plot(figsize=(12, 8), legend=True, label='Exisiting Points')
         future_preds.plot(figsize=(12, 8), legend=True, label='Future Points')
```

```
Out[17]: <Axes: xlabel='Order Date'>
```




```
In [18]: df = pd.read_csv('C://Users//saswa//OneDrive//Desktop//Pinaki-Time-series-forecasting//Pinaki-samples.csv', index_col=0, parse_dates=True)
df.head()
```

Out[18]:


	a	b	c	d
1950-01-01	36	27	0	67
1950-02-01	58	22	3	31
1950-03-01	61	17	5	67
1950-04-01	37	15	8	47
1950-05-01	66	13	8	62

```
In [19]: df['a'].plot(ylim=[0, 100], title='STATIONARY DATA').autoscale(axis='x', tight=True)
```




```
In [20]: df['a_ma'] = df['a'].rolling(12).mean()  
  
df['a_ma'].plot(ylim=[0, 100], title='Moving Average')
```

```
Out[20]: <Axes: title={'center': 'Moving Average'}>
```




```
In [21]: df['b'].plot(ylim=[0,100],title="NON-STATIONARY DATA").autoscale(axis='x',tight=True)
```




```
In [22]: df['b_ma'] = df['b'].rolling(12).mean()
df['b_ma'].plot(ylim=[0, 100], title='Moving Average')
```

```
Out[22]: <Axes: title={'center': 'Moving Average'}>
```




```
In [23]: from statsmodels.tsa.statespace.tools import diff  
  
df['d1'] = diff(df['b'],k_diff=1)  
  
df['d1'].plot(title="FIRST DIFFERENCE DATA", ylim=[-30, 30]).autoscale(axis='x',tight=True)
```



```
In [24]: df['d1_ma'] = df['d1'].rolling(12).mean()

df['d1_ma'].plot(title="MA OF FIRST DIFFERENCE DATA", ylim=[-10, 10]).autoscale(axis='x',tight=True)
```




```
In [25]: import pandas as pd
import numpy as np
%matplotlib inline
import statsmodels.api as sm
```

```
In [26]: import warnings
warnings.filterwarnings("ignore")
```

```
In [27]: df = pd.DataFrame({'a':[1, 10, 7, 13, 6, 18, 9, 25, 13]})

df.plot()
```

```
Out[27]: <Axes: >
```

```
In [28]: from statsmodels.tsa.stattools import acovf, acf
arr = acovf(df['a'])

arr
```

```
Out[28]: array([ 44.22222222, -6.30864198, 19.60493827, -11.62962963,
        7.80246914, -13.58024691, -0.14814815, -15.9382716 ,
        -1.91358025])
```

```
In [29]: arr1 = acovf(df['a'])

arr1
```

```
Out[29]: array([ 44.22222222, -6.30864198,  19.60493827, -11.62962963,
                7.80246914, -13.58024691, -0.14814815, -15.9382716 ,
               -1.91358025])
```

```
In [30]: arr3 = acf(df['a'])
```

```
arr3
```

```
Out[30]: array([ 1.          , -0.14265773,  0.44332775, -0.26298157,  0.17643774,
               -0.30709101, -0.00335008, -0.36041318, -0.04327192])
```

```
In [31]: from statsmodels.tsa.stattools import pacf_yw
arr4 = pacf_yw(df['a'],nlags=4,method='mle')
arr4
```

```
Out[31]: array([ 1.          , -0.14265773,  0.43176344, -0.20758442, -0.04572862])
```

```
In [32]: arr4 = pacf_yw(df['a'],nlags=4,method='adjusted')
arr4
```

```
Out[32]: array([ 1.          , -0.16048995,  0.5586243 , -0.39456104,  0.01906252])
```

```
In [33]: from statsmodels.tsa.stattools import pacf_ols
arr5 = pacf_ols(df['a'],nlags=4)
arr5
```

```
Out[33]: array([ 1.          , -0.13833492,  1.13495418, -0.04476691,  0.5979815 ])
```

```
In [34]: from pandas.plotting import lag_plot
```

```
In [36]: df = pd.read_excel('C://Users//saswa//OneDrive//Desktop//Pinaki-Time-series-forecasting//Superstore_Sales_Records.xls', index_
df['Category'].value_counts())
```

```
Out[36]: Office Supplies    6026
Furniture                  2121
Technology                 1847
Name: Category, dtype: int64
```

```
In [37]: df = pd.read_excel('C://Users//saswa//OneDrive//Desktop//Pinaki-Time-series-forecasting//Superstore_Sales_Records.xls', index_
df = df[df['Category']=='Office Supplies']
```

```
df = df.groupby(by='Order Date').agg({'Sales':sum})  
df.sort_index(inplace=True)  
df.head(4)
```

Out[37]:

	Sales
--	-------

Order Date	
------------	--

2014-01-03	16.448
------------	--------

2014-01-04	288.060
------------	---------

2014-01-05	19.536
------------	--------

2014-01-06	685.340
------------	---------

```
In [38]: df = df.resample('W').sum()  
  
df.head()
```

Out[38]:

	Sales
--	-------

Order Date	
------------	--

2014-01-05	324.044
------------	---------

2014-01-12	708.004
------------	---------


2014-01-19	2337.764
------------	----------

2014-01-26	1143.170
------------	----------

2014-02-02	368.784
------------	---------

```
In [39]: lag_plot(df['Sales'])
```

Out[39]: <Axes: xlabel='y(t)', ylabel='y(t + 1)'>





```
In [40]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
         acf(df['Sales'])
```

```
Out[40]: array([ 1.00000000e+00,  3.12440082e-01,  2.49173025e-01,  1.39336867e-01,
                  1.84489786e-01,  1.67531418e-01,  1.11379734e-01,  1.84930768e-01,
                  8.24198806e-02,  8.74684994e-02,  4.89924777e-02,  6.22098365e-02,
                  1.65448013e-01,  1.49072360e-01,  7.97573259e-02,  7.65911202e-02,
                  7.55594141e-02, -7.46386648e-04,  4.97054222e-02, -2.10402529e-02,
                  -2.49880823e-02, -6.68093251e-02, -1.79718855e-02, -3.42941758e-02])
```

```
In [41]: title = 'Autocorrelation: Weekly Sales of Office Supplies'
         lags = 40
         plot_acf(df['Sales'], title=title)
```

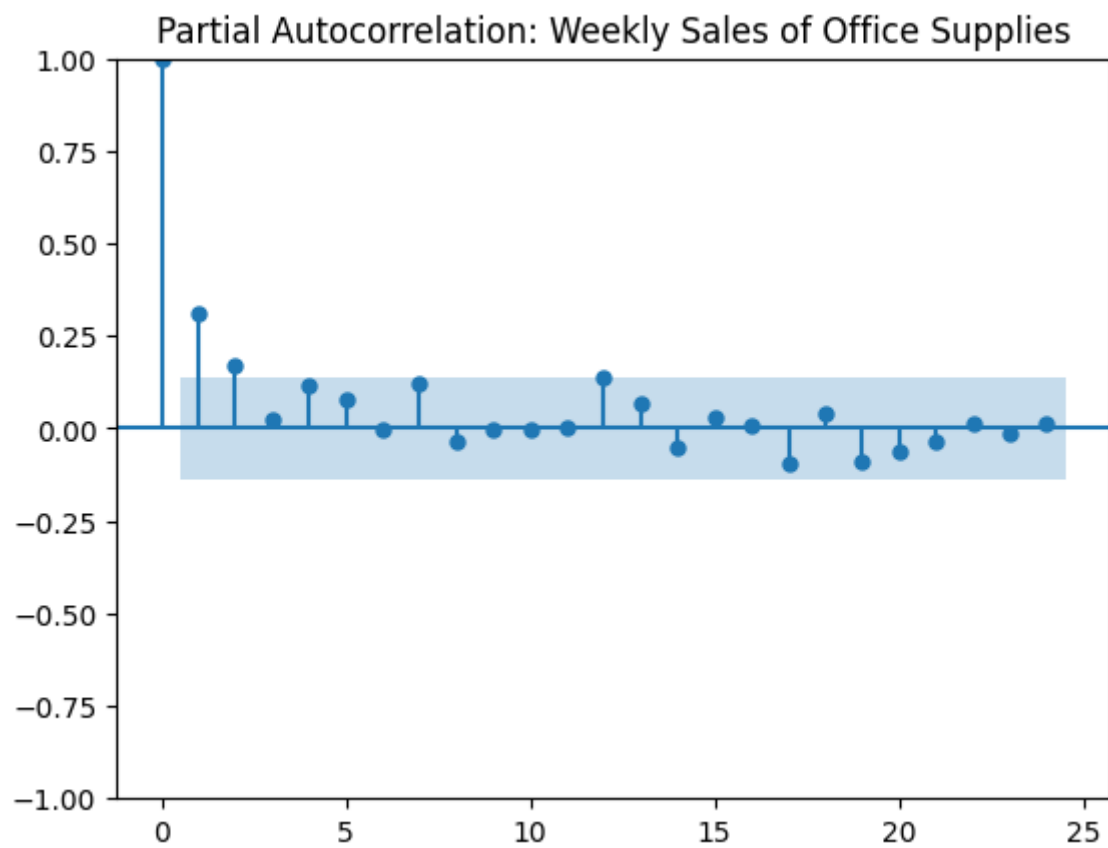
Out[41]:






```
In [42]: title = 'Partial Autocorrelation: Weekly Sales of Office Supplies'
lags = 40
plot_pacf(df['Sales'], title=title)
```

Out[42]:





```
In [43]: import pandas as pd
import numpy as np
%matplotlib inline
```

```
In [44]: from statsmodels.tsa.ar_model import AR,ARResults
```

```
#plot the sales data of Office Supplies
```

```
df = pd.read_excel('C://Users//saswa//OneDrive//Desktop//Pinaki-Time-series-forecasting//Superstore_Sales_Records.xls', index_
df = df[df['Category']=='Furniture']
df = df.groupby(by='Order Date').agg({'Sales':sum})
df.sort_index(inplace=True)
df.head(4)
```


Out[44]:

Sales	
Order Date	
2014-01-06	2573.820
2014-01-07	76.728
2014-01-10	51.940
2014-01-11	9.940

```
In [45]: df = df.resample('M').sum()
```

```
df.head()
```

Out[45]:

Sales	
Order Date	
2014-01-31	6242.525
2014-02-28	1839.658
2014-03-31	14573.956
2014-04-30	7944.837
2014-05-31	6912.787

```
In [46]: title='Monthly Furniture Sales Data'
```

```
ylabel='Sales Data'
```


```
xlabel='Order Date'
```

```
ax = df['Sales'].plot(figsize=(12,5),title=title)
```

```
ax.autoscale(axis='x',tight=True)
```

```
ax.set(xlabel=xlabel, ylabel=ylabel)
```

```
Out[46]: [Text(0.5, 0, 'Order Date'), Text(0, 0.5, 'Sales Data')]
```



```
In [47]: len(df)
```

```
Out[47]: 48
```

```
In [48]: train = df.iloc[:len(df)-6]
test = df.iloc[len(df)-6:]

print(f"Train size is {len(train)}")
print(f"Test size is {len(test)}")
```

```
Train size is 42
```

```
Test size is 6
```

```
In [49]: import warnings
```

```
warnings.filterwarnings("ignore")

from statsmodels.tsa.ar_model import AutoReg
```

```
In [50]: mod1 = AutoReg(train['Sales'], 1, old_names=False)
res1 = mod1.fit()
print(res1.summary())
```

```

                        AutoReg Model Results
=====
Dep. Variable:            Sales   No. Observations:                42
Model:                    AutoReg(1)   Log Likelihood                -427.247
Method:                    Conditional MLE   S.D. of innovations            8116.944
Date:                      Thu, 21 Sep 2023   AIC                           860.493
Time:                      00:07:29   BIC                           865.634
Sample:                    02-28-2014   HQIC                          862.365
                        - 06-30-2017
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
const      1.172e+04   2487.380      4.712      0.000     6845.693     1.66e+04
Sales.L1    0.1884      0.152      1.237      0.216     -0.110      0.487
=====
                        Roots
=====
              Real      Imaginary      Modulus      Frequency
-----
AR.1           5.3086      +0.0000j           5.3086           0.0000
=====
```

```
In [51]: print(f'Lag: {res1.arfreq}')
print(f'Coefficients:\n{res1.params}')
```

```
Lag: [0.]
Coefficients:
const      11720.868755
Sales.L1    0.188372
dtype: float64
```

```
In [52]: start=len(train)
end=len(train)+len(test)-1
predictions1 = res1.predict(start=start, end=end, dynamic=False).rename('AR(1) Predictions')
```

```
In [53]: predictions1
```


```
Out[53]: 2017-07-31    15301.562179
          2017-08-31    14603.261009
          2017-09-30    14471.720330
          2017-10-31    14446.941694
          2017-11-30    14442.274083
          2017-12-31    14441.394833
          Freq: M, Name: AR(1) Predictions, dtype: float64
```

```
In [54]: for i in range(len(predictions1)):
          print(f"predicted={predictions1[i]:<11.10}, expected={test['Sales'][i]}")
```

```
predicted=15301.56218, expected=11813.021999999999
predicted=14603.26101, expected=15441.874
predicted=14471.72033, expected=29028.206000000002
predicted=14446.94169, expected=21884.0682
predicted=14442.27408, expected=37056.715
predicted=14441.39483, expected=31407.4668
```

```
In [55]: test['Sales'].plot(legend=True)
          predictions1.plot(legend=True,figsize=(12,6))
```

```
Out[55]: <Axes: xlabel='Order Date'>
```



```
In [56]: mod6 = AutoReg(train['Sales'], 6)
res6 = mod6.fit()
print(res6.summary())
```

```

AutoReg Model Results
=====
Dep. Variable:      Sales    No. Observations:      42
Model:              AutoReg(6)  Log Likelihood          -373.443
Method:             Conditional MLE  S.D. of innovations      7742.555
Date:              Thu, 21 Sep 2023  AIC              762.887
Time:              00:08:29    BIC              775.555
Sample:            07-31-2014    HQIC             767.308
                  - 06-30-2017

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
const      1.93e+04    5788.110      3.334      0.001     7950.646     3.06e+04
Sales.L1      0.1277      0.165      0.774      0.439      -0.196      0.451
Sales.L2     -0.0805      0.162     -0.496      0.620      -0.399      0.238
Sales.L3      0.0557      0.159      0.350      0.726      -0.256      0.367
Sales.L4     -0.2548      0.158     -1.610      0.107      -0.565      0.055
Sales.L5     -0.1875      0.160     -1.175      0.240      -0.500      0.125
Sales.L6      0.0541      0.160      0.339      0.735      -0.259      0.367

Roots
=====
              Real      Imaginary      Modulus      Frequency
-----
AR.1      -0.5872      -1.1332j      1.2763      -0.3261
AR.2      -0.5872      +1.1332j      1.2763      0.3261
AR.3      -1.6817      -0.0000j      1.6817      -0.5000
AR.4       0.9181      -0.8134j      1.2266      -0.1154
AR.5       0.9181      +0.8134j      1.2266      0.1154
AR.6       4.4871      -0.0000j      4.4871      -0.0000
=====

```

```


In [57]: start=len(train)
end=len(train)+len(test)-1
predictions6 = res6.predict(start=start, end=end, dynamic=False).rename('AR(6) Predictions')
predictions6

```

```
Out[57]: 2017-07-31    17121.629045  
         2017-08-31    16914.370614  
         2017-09-30    15703.623235  
         2017-10-31    13359.612451  
         2017-11-30    13669.199630  
         2017-12-31    14347.505858  
         Freq: M, Name: AR(6) Predictions, dtype: float64
```

```
In [58]: test['Sales'].plot(legend=True)  
         predictions6.plot(legend=True,figsize=(12,6))
```

```
Out[58]: <Axes: xlabel='Order Date'>
```



```
In [59]: from statsmodels.tsa.ar_model import ar_select_order
p = ar_select_order(train['Sales'], maxlag=15)
p.ar_lags
```

```
Out[59]: [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12]
```

```
In [60]: mod12 = AutoReg(train['Sales'], 12)
res12 = mod12.fit()
print(res12.summary())
```



AutoReg Model Results						
=====						
Dep. Variable:	Sales	No. Observations:	42			
Model:	AutoReg(12)	Log Likelihood	-290.930			
Method:	Conditional MLE	S.D. of innovations	3939.183			
Date:	Thu, 21 Sep 2023	AIC	609.860			
Time:	00:09:12	BIC	629.477			
Sample:	01-31-2015	HQIC	616.136			
	- 06-30-2017					
=====						
	coef	std err	z	P> z	[0.025	0.975]

const	6179.8244	9948.341	0.621	0.534	-1.33e+04	2.57e+04
Sales.L1	0.0175	0.102	0.172	0.864	-0.182	0.217
Sales.L2	-0.0058	0.104	-0.056	0.956	-0.209	0.197
Sales.L3	-0.0491	0.102	-0.481	0.631	-0.249	0.151
Sales.L4	-0.1194	0.103	-1.163	0.245	-0.320	0.082
Sales.L5	-0.0596	0.101	-0.593	0.553	-0.257	0.138
Sales.L6	-0.0091	0.101	-0.091	0.928	-0.206	0.188
Sales.L7	0.0419	0.114	0.369	0.712	-0.181	0.265
Sales.L8	0.0262	0.118	0.221	0.825	-0.206	0.258
Sales.L9	-0.0595	0.115	-0.516	0.606	-0.285	0.166
Sales.L10	-0.0724	0.123	-0.590	0.555	-0.313	0.168
Sales.L11	0.0335	0.118	0.284	0.776	-0.198	0.265
Sales.L12	0.9403	0.115	8.143	0.000	0.714	1.167
Roots						
=====						
	Real	Imaginary	Modulus	Frequency		

AR.1	1.0270	-0.0000j	1.0270	-0.0000		
AR.2	0.8601	-0.5085j	0.9991	-0.0850		
AR.3	0.8601	+0.5085j	0.9991	0.0850		
AR.4	0.5128	-0.8447j	0.9882	-0.1632		
AR.5	0.5128	+0.8447j	0.9882	0.1632		
AR.6	-0.0108	-1.0048j	1.0049	-0.2517		
AR.7	-0.0108	+1.0048j	1.0049	0.2517		
AR.8	-0.5131	-0.8690j	1.0091	-0.3349		
AR.9	-0.5131	+0.8690j	1.0091	0.3349		
AR.10	-0.8731	-0.5059j	1.0091	-0.4164		
AR.11	-0.8731	+0.5059j	1.0091	0.4164		
AR.12	-1.0144	-0.0000j	1.0144	-0.5000		

```
In [61]: start=len(train)
end=len(train)+len(test)-1
predictions12 = res12.predict(start=start, end=end, dynamic=False).rename('AR(6) Predictions')
test['Sales'].plot(legend=True)
predictions12.plot(legend=True,figsize=(12,6))
```

Out[61]: <Axes: xlabel='Order Date'>



```
In [62]: from sklearn.metrics import mean_squared_error
```

```
labels = ['AR(1)', 'AR(6)', 'AR(12)']
preds = [predictions1, predictions6, predictions12] # these are variables, not strings!

for i in range(3):
    error = mean_squared_error(test['Sales'], preds[i])
    print(f'{labels[i]} Error: {error:11.10}')
```

AR(1) Error: 179889307.7

AR(6) Error: 186429758.2

AR(12) Error: 30320630.65

In [63]: modls = [res1, res6, res12]

```
for i in range(3):
    print(f'{labels[i]} AIC: {modls[i].aic:6.5}')
```

AR(1) AIC: 860.49

AR(6) AIC: 762.89

AR(12) AIC: 609.86

In [64]: mod12 = AutoReg(df['Sales'], 12)
res12 = mod12.fit()
print(res12.summary())

AutoReg Model Results						
=====						
Dep. Variable:	Sales	No. Observations:	48			
Model:	AutoReg(12)	Log Likelihood	-349.076			
Method:	Conditional MLE	S.D. of innovations	3934.790			
Date:	Thu, 21 Sep 2023	AIC	726.152			
Time:	00:10:01	BIC	748.321			
Sample:	01-31-2015	HQIC	733.889			
	- 12-31-2017					
=====						
	coef	std err	z	P> z	[0.025	0.975]

const	4356.0138	7562.612	0.576	0.565	-1.05e+04	1.92e+04
Sales.L1	0.0321	0.088	0.366	0.715	-0.140	0.204
Sales.L2	-0.0097	0.094	-0.102	0.918	-0.195	0.175
Sales.L3	-0.0670	0.092	-0.727	0.467	-0.248	0.114
Sales.L4	-0.0820	0.097	-0.846	0.398	-0.272	0.108
Sales.L5	-0.0348	0.094	-0.369	0.712	-0.219	0.150
Sales.L6	0.0304	0.095	0.320	0.749	-0.155	0.216
Sales.L7	-0.0261	0.093	-0.279	0.780	-0.209	0.157
Sales.L8	-0.0317	0.093	-0.339	0.734	-0.215	0.151
Sales.L9	-0.0284	0.094	-0.303	0.762	-0.212	0.155
Sales.L10	0.0041	0.093	0.044	0.965	-0.178	0.186
Sales.L11	0.1429	0.090	1.592	0.111	-0.033	0.319
Sales.L12	0.8895	0.089	9.993	0.000	0.715	1.064
Roots						
=====						
	Real	Imaginary	Modulus	Frequency		

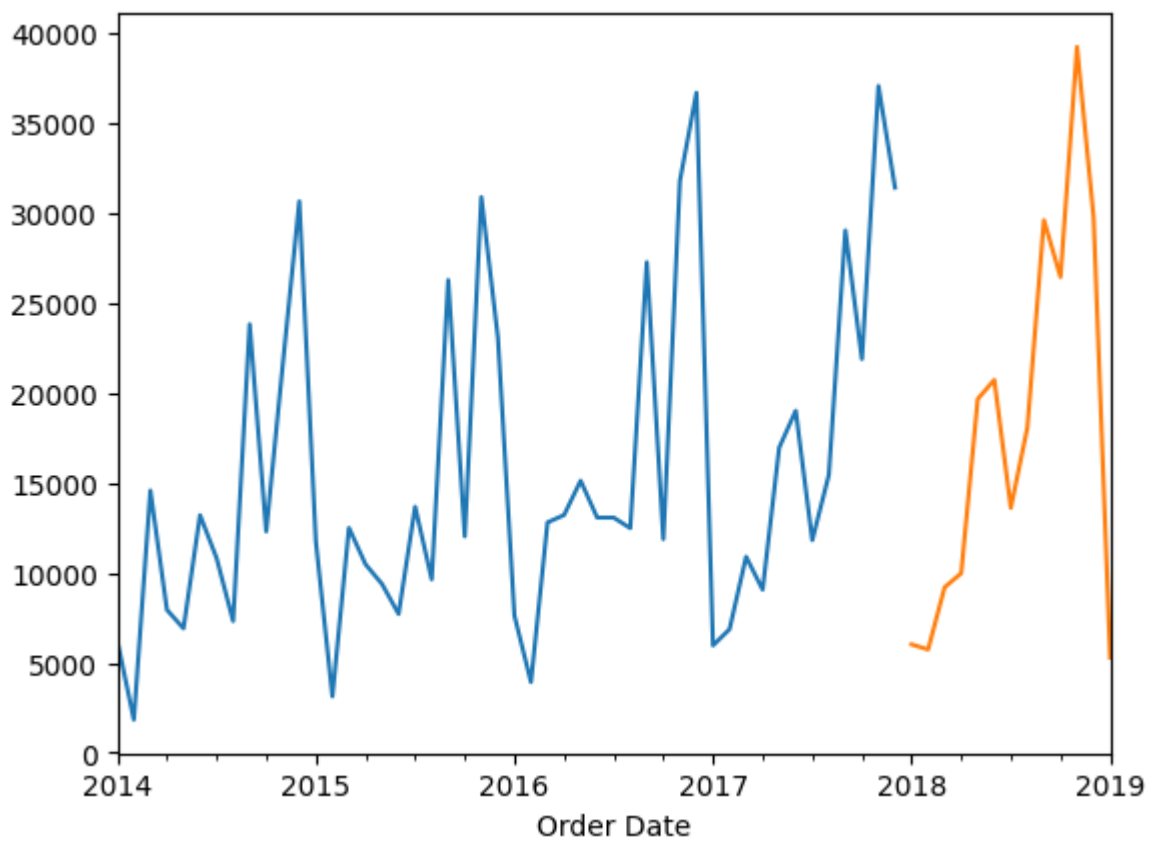
AR.1	1.0149	-0.0000j	1.0149	-0.0000		
AR.2	0.8510	-0.5063j	0.9902	-0.0854		
AR.3	0.8510	+0.5063j	0.9902	0.0854		
AR.4	0.4907	-0.8578j	0.9882	-0.1673		
AR.5	0.4907	+0.8578j	0.9882	0.1673		
AR.6	-0.0094	-1.0212j	1.0212	-0.2515		
AR.7	-0.0094	+1.0212j	1.0212	0.2515		
AR.8	-0.5178	-0.8730j	1.0150	-0.3352		
AR.9	-0.5178	+0.8730j	1.0150	0.3352		
AR.10	-1.0205	-0.0000j	1.0205	-0.5000		
AR.11	-0.8920	-0.5093j	1.0272	-0.4174		
AR.12	-0.8920	+0.5093j	1.0272	0.4174		

```
-----  
In [65]: start = len(df)  
end = len(df)+12  
  
pred_future = res12.predict(start=start, end=end, dynamic=False)  
  
pred_future
```

```
Out[65]: 2018-01-31    6022.626972  
2018-02-28    5739.250407  
2018-03-31    9199.016790  
2018-04-30    9976.759362  
2018-05-31   19661.613494  
2018-06-30   20722.318428  
2018-07-31   13609.961837  
2018-08-31   18080.731150  
2018-09-30   29604.129926  
2018-10-31   26437.359412  
2018-11-30   39227.722049  
2018-12-31   29781.186936  
2019-01-31    5288.024913  
Freq: M, dtype: float64
```

```
In [66]: df['Sales'].plot()  
pred_future.plot()
```

```
Out[66]: <Axes: xlabel='Order Date'>
```



```
In [67]: pip install pmdarima
```

```

Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: pmdarima in c:\users\saswa\appdata\roaming\python\python311\site-packages (2.0.3)
Requirement already satisfied: joblib>=0.11 in c:\users\saswa\appdata\roaming\python\python311\site-packages (from pmdarima) (1.3.1)
Requirement already satisfied: Cython!=0.29.18,!0.29.31,>=0.29 in c:\users\saswa\appdata\roaming\python\python311\site-packages (from pmdarima) (3.0.2)
Requirement already satisfied: numpy>=1.21.2 in c:\users\saswa\appdata\roaming\python\python311\site-packages (from pmdarima) (1.24.2)
Requirement already satisfied: pandas>=0.19 in c:\users\saswa\appdata\roaming\python\python311\site-packages (from pmdarima) (1.5.3)
Requirement already satisfied: scikit-learn>=0.22 in c:\users\saswa\appdata\roaming\python\python311\site-packages (from pmdarima) (1.3.0)
Requirement already satisfied: scipy>=1.3.2 in c:\users\saswa\appdata\roaming\python\python311\site-packages (from pmdarima) (1.11.2)
Requirement already satisfied: statsmodels>=0.13.2 in c:\users\saswa\appdata\roaming\python\python311\site-packages (from pmdarima) (0.14.0)
Requirement already satisfied: urllib3 in c:\users\saswa\appdata\roaming\python\python311\site-packages (from pmdarima) (2.0.4)
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in c:\program files\python311\lib\site-packages (from pmdarima) (65.5.0)
Requirement already satisfied: python-dateutil>=2.8.1 in c:\users\saswa\appdata\roaming\python\python311\site-packages (from pandas>=0.19->pmdarima) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in c:\users\saswa\appdata\roaming\python\python311\site-packages (from pandas>=0.19->pmdarima) (2022.7.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\saswa\appdata\roaming\python\python311\site-packages (from scikit-learn>=0.22->pmdarima) (3.2.0)
Requirement already satisfied: patsy>=0.5.2 in c:\users\saswa\appdata\roaming\python\python311\site-packages (from statsmodels>=0.13.2->pmdarima) (0.5.3)
Requirement already satisfied: packaging>=21.3 in c:\users\saswa\appdata\roaming\python\python311\site-packages (from statsmodels>=0.13.2->pmdarima) (23.1)
Requirement already satisfied: six in c:\users\saswa\appdata\roaming\python\python311\site-packages (from patsy>=0.5.2->statsmodels>=0.13.2->pmdarima) (1.16.0)
Note: you may need to restart the kernel to use updated packages.

```

```

In [68]: import pandas as pd
import numpy as np
%matplotlib inline

# Load specific forecasting tools
from statsmodels.tsa.arima_model import ARMA, ARMAResults, ARIMA, ARIMAResults
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf # for determining (p,q) orders
from pmdarima import auto_arima # for determining ARIMA orders

```

```
# Ignore harmless warnings
import warnings
warnings.filterwarnings("ignore")
```

```
In [69]: df = pd.read_excel('C://Users//saswa//OneDrive//Desktop//Pinaki-Time-series-forecasting//Superstore_Sales_Records.xls', index_
df = df[df['Category']=='Furniture']
df = df.groupby(by='Order Date').agg({'Sales':sum})
df.sort_index(inplace=True)
df.head(4)
```

Out[69]:

	Sales
--	-------

Order Date	
------------	--

2014-01-06	2573.820
------------	----------

2014-01-07	76.728
------------	--------

2014-01-10	51.940
------------	--------

2014-01-11	9.940
------------	-------

```
In [70]: df = df.resample('MS').sum()
df.head()
```

Out[70]:

	Sales
--	-------

Order Date	
------------	--

2014-01-01	6242.525
------------	----------

2014-02-01	1839.658
------------	----------

2014-03-01	14573.956
------------	-----------

2014-04-01	7944.837
------------	----------

2014-05-01	6912.787
------------	----------


```
In [71]: from statsmodels.tsa.stattools import adfuller
```

```
In [72]: def adf_test(series,title=''):
    """
    Pass in a time series and an optional title, returns an ADF report
    """
    print(f'Augmented Dickey-Fuller Test: {title}')
    result = adfuller(series.dropna(),autolag='AIC')

    labels = ['ADF test statistic','p-value','# lags used','# observations']
    out = pd.Series(result[0:4],index=labels)


    for key,val in result[4].items():
        out[f'critical value ({key})']=val

    print(out.to_string())

    if result[1] <= 0.05:
        print("Strong evidence against the null hypothesis")
        print("Reject the null hypothesis")
        print("Data has no unit root and is stationary")
    else:
        print("Weak evidence against the null hypothesis")
        print("Fail to reject the null hypothesis")
        print("Data has a unit root and is non-stationary")
```

```
In [73]: df['Sales'].plot()
```

```
Out[73]: <Axes: xlabel='Order Date'>
```



```
In [74]: adf_test(df['Sales'])
```

Augmented Dickey-Fuller Test:

ADF test statistic -4.699026

p-value 0.000085

lags used 0.000000

observations 47.000000

critical value (1%) -3.577848

critical value (5%) -2.925338

critical value (10%) -2.600774

Strong evidence against the null hypothesis

Reject the null hypothesis

Data has no unit root and is stationary

```
In [75]: auto_arima(df['Sales'], seasonal=True).summary()
```

```
Out[75]:
```

SARIMAX Results

Dep. Variable:	y	No. Observations:	48			
Model:	SARIMAX(1, 0, 0)	Log Likelihood	-502.820			
Date:	Thu, 21 Sep 2023	AIC	1011.640			
Time:	00:14:34	BIC	1017.253			
Sample:	01-01-2014	HQIC	1013.761			
	- 12-01-2017					
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
intercept	1.084e+04	2695.066	4.021	0.000	5554.237	1.61e+04
ar.L1	0.3056	0.131	2.328	0.020	0.048	0.563
sigma2	7.318e+07	0.160	4.56e+08	0.000	7.32e+07	7.32e+07
Ljung-Box (L1) (Q):	0.00	Jarque-Bera (JB):	3.70			
Prob(Q):	0.98	Prob(JB):	0.16			
Heteroskedasticity (H):	1.88	Skew:	0.64			
Prob(H) (two-sided):	0.22	Kurtosis:	2.54			

Warnings:

- [1] Covariance matrix calculated using the outer product of gradients (complex-step).
- [2] Covariance matrix is singular or near-singular, with condition number 2.09e+24. Standard errors may be unstable.

```
In [76]: train = df.iloc[:len(df)-6]
        test = df.iloc[len(df)-6:]

        len(train), len(test)
```

```
Out[76]: (42, 6)
```

```
In [77]: help(ARMA)
```

Help on class ARMA in module statsmodels.tsa.arima_model:

```
class ARMA(builtins.object)
|   ARMA(*args, **kwargs)
|
|   ARMA has been deprecated in favor of the new implementation
|
|   See Also
|   -----
|   statsmodels.tsa.arima.model.ARIMA
|       ARIMA models with a variety of parameter estimators
|   statsmodels.tsa.statespace.SARIMAX
|       SARIMAX models estimated using MLE
|
|   Methods defined here:
|
|   __init__(self, *args, **kwargs)
|       Initialize self.  See help(type(self)) for accurate signature.
|
|   -----
|   Data descriptors defined here:
|
|   __dict__
|       dictionary for instance variables (if defined)
|
|   __weakref__
|       list of weak references to the object (if defined)
```

```
In [78]: from statsmodels.tsa.arima.model import ARIMA
        model = ARIMA(train['Sales'], order=(2,0,2))
        results = model.fit()
```

```
results.summary()
```

Out[78]:

SARIMAX Results

Dep. Variable:		Sales	No. Observations:		42	
Model:		ARIMA(2, 0, 2)		Log Likelihood		-436.980
Date:		Thu, 21 Sep 2023		AIC		885.960
Time:		00:15:08		BIC		896.386
Sample:		01-01-2014		HQIC		889.782
		- 06-01-2017				
Covariance Type:		opg				
	coef	std err	z	P> z	[0.025	0.975]
const	1.418e+04	1427.017	9.934	0.000	1.14e+04	1.7e+04
ar.L1	0.6732	0.810	0.831	0.406	-0.914	2.260
ar.L2	0.1773	0.842	0.211	0.833	-1.474	1.828
ma.L1	-0.5278	0.739	-0.714	0.475	-1.977	0.921
ma.L2	-0.4528	0.771	-0.587	0.557	-1.964	1.059
sigma2	7.04e+07	0.014	4.87e+09	0.000	7.04e+07	7.04e+07
Ljung-Box (L1) (Q):		0.01	Jarque-Bera (JB):		8.53	
Prob(Q):		0.93	Prob(JB):		0.01	
Heteroskedasticity (H):		1.31	Skew:		1.10	
Prob(H) (two-sided):		0.62	Kurtosis:		3.27	

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 2.64e+25. Standard errors may be unstable.


```
In [79]: start=len(train)
end=len(train)+len(test)-1
predictions = results.predict(start=start, end=end).rename('ARMA(2,2) Predictions')
predictions
```

```
Out[79]: 2017-07-01    13115.074287
2017-08-01    12031.112128
2017-09-01    12543.910534
2017-10-01    12696.894212
2017-11-01    12890.811988
2017-12-01    13048.480070
Freq: MS, Name: ARMA(2,2) Predictions, dtype: float64
```

```
In [80]: title = 'Furniture Sales Data'
ylabel='Sales Data'
xlabel='Order Date'

ax = test['Sales'].plot(legend=True,figsize=(12,6),title=title)
predictions.plot(legend=True)
ax.autoscale(axis='x',tight=True)
ax.set(xlabel=xlabel, ylabel=ylabel)
```

```
Out[80]: [Text(0.5, 0, 'Order Date'), Text(0, 0.5, 'Sales Data')]
```



```
In [81]: df = pd.read_excel('C://Users//saswa//OneDrive//Desktop//Pinaki-Time-series-forecasting//Superstore_Sales_Records.xls', index_
df = df[df['Category']=='Furniture']
df = df.groupby(by='Order Date').agg({'Sales':sum})
df.sort_index(inplace=True)
df.head(4)
```


Out[81]:

	Sales
Order Date	
2014-01-06	2573.820
2014-01-07	76.728
2014-01-10	51.940
2014-01-11	9.940

```
In [82]: df = df.resample('MS').sum()
df.head()
```


Out[82]:

	Sales
Order Date	
2014-01-01	6242.525
2014-02-01	1839.658
2014-03-01	14573.956
2014-04-01	7944.837
2014-05-01	6912.787

```
In [83]: import matplotlib.ticker as ticker
formatter = ticker.StrMethodFormatter('{x:,.0f}')

title = 'Furniture Sales Data'
ylabel='Sales Data'
xlabel='Order Date'


ax = df['Sales'].plot(figsize=(12,5),title=title)
ax.autoscale(axis='x',tight=True)
ax.set(xlabel=xlabel, ylabel=ylabel)
ax.yaxis.set_major_formatter(formatter);
```




```
In [84]: from statsmodels.tsa.seasonal import seasonal_decompose

result = seasonal_decompose(df['Sales'], model='additive')
result.plot()
```

Out[84]:





```
In [85]: auto_arima(df['Sales'], seasonal=False).summary()
```

Out[85]:

SARIMAX Results

Dep. Variable:		y	No. Observations:		48	
Model:		SARIMAX(1, 0, 0)		Log Likelihood		-502.820
Date:		Thu, 21 Sep 2023		AIC		1011.640
Time:		00:17:17		BIC		1017.253
Sample:		01-01-2014		HQIC		1013.761
		- 12-01-2017				
Covariance Type:		opg				
	coef	std err	z	P> z	[0.025	0.975]
intercept	1.084e+04	2695.066	4.021	0.000	5554.237	1.61e+04
ar.L1	0.3056	0.131	2.328	0.020	0.048	0.563
sigma2	7.318e+07	0.160	4.56e+08	0.000	7.32e+07	7.32e+07
Ljung-Box (L1) (Q):		0.00	Jarque-Bera (JB):		3.70	
Prob(Q):		0.98	Prob(JB):		0.16	
Heteroskedasticity (H):		1.88	Skew:		0.64	
Prob(H) (two-sided):		0.22	Kurtosis:		2.54	

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 2.09e+24. Standard errors may be unstable.

In [86]: `from statsmodels.tsa.statespace.tools import diff`

```
df['d1'] = diff(df['Sales'], k_diff=1)

adf_test(df['d1'], 'Furniture Sales Data')
```

Augmented Dickey-Fuller Test: Furniture Sales Data

ADF test statistic -1.147459e+01

p-value 5.167971e-21

lags used 1.000000e+01

observations 3.600000e+01

critical value (1%) -3.626652e+00

critical value (5%) -2.945951e+00

critical value (10%) -2.611671e+00


Strong evidence against the null hypothesis


Reject the null hypothesis

Data has no unit root and is stationary

```
In [87]: title = 'Furniture Monthly Sales Data'
         lags = 40
         plot_acf(df['Sales'],title=title,lags=lags)
```


Out[87]:






```
In [88]: title = 'Partial Autocorrelation: Monthly Furniture Sales Data'
lags = 12
plot_pacf(df['Sales'],title=title,lags=lags)
```


Out[88]:





```
In [89]: stepwise_fit = auto_arima(df['Sales'], start_p=0, start_q=0,
                                max_p=2, max_q=2, m=12,
                                seasonal=False,
                                d=None, trace=True,
                                error_action='ignore', # we don't want to know if an order does not work
                                suppress_warnings=True, # we don't want convergence warnings
                                stepwise=True)         # set to stepwise

stepwise_fit.summary()
```

Performing stepwise search to minimize aic

ARIMA(0,0,0)(0,0,0)[0]	: AIC=1078.259, Time=0.02 sec
ARIMA(1,0,0)(0,0,0)[0]	: AIC=1026.129, Time=0.04 sec
ARIMA(0,0,1)(0,0,0)[0]	: AIC=1059.367, Time=0.05 sec
ARIMA(2,0,0)(0,0,0)[0]	: AIC=1022.619, Time=0.06 sec
ARIMA(2,0,1)(0,0,0)[0]	: AIC=1018.040, Time=0.18 sec
ARIMA(1,0,1)(0,0,0)[0]	: AIC=1018.347, Time=0.14 sec
ARIMA(2,0,2)(0,0,0)[0]	: AIC=1020.706, Time=0.44 sec
ARIMA(1,0,2)(0,0,0)[0]	: AIC=1019.172, Time=0.15 sec
ARIMA(2,0,1)(0,0,0)[0] intercept	: AIC=1015.561, Time=0.15 sec
ARIMA(1,0,1)(0,0,0)[0] intercept	: AIC=1013.580, Time=0.08 sec
ARIMA(0,0,1)(0,0,0)[0] intercept	: AIC=1011.909, Time=0.15 sec
ARIMA(0,0,0)(0,0,0)[0] intercept	: AIC=1014.386, Time=0.08 sec
ARIMA(0,0,2)(0,0,0)[0] intercept	: AIC=1013.938, Time=0.09 sec
ARIMA(1,0,0)(0,0,0)[0] intercept	: AIC=1011.640, Time=0.06 sec
ARIMA(2,0,0)(0,0,0)[0] intercept	: AIC=1013.680, Time=0.18 sec

Best model: ARIMA(1,0,0)(0,0,0)[0] intercept

Total fit time: 1.890 seconds

Out[89]:

SARIMAX Results

Dep. Variable:		y	No. Observations:		48	
Model:		SARIMAX(1, 0, 0)		Log Likelihood		-502.820
Date:		Thu, 21 Sep 2023		AIC		1011.640
Time:		00:18:13		BIC		1017.253
Sample:		01-01-2014		HQIC		1013.761
		- 12-01-2017				
Covariance Type:		opg				
	coef	std err	z	P> z	[0.025	0.975]
intercept	1.084e+04	2695.066	4.021	0.000	5554.237	1.61e+04
ar.L1	0.3056	0.131	2.328	0.020	0.048	0.563
sigma2	7.318e+07	0.160	4.56e+08	0.000	7.32e+07	7.32e+07
Ljung-Box (L1) (Q):		0.00	Jarque-Bera (JB):		3.70	
Prob(Q):		0.98	Prob(JB):		0.16	
Heteroskedasticity (H):		1.88	Skew:		0.64	
Prob(H) (two-sided):		0.22	Kurtosis:		2.54	

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 2.09e+24. Standard errors may be unstable.

```
In [90]: n = len(df)
test_size = 12
```

```
train = df.iloc[:n-test_size]  
test = df.iloc[n-test_size:]
```

```
In [91]: model = ARIMA(train['Sales'], order=(1, 0, 0))  
  
results = model.fit()  
results.summary()
```

Out[91]:

SARIMAX Results

Dep. Variable:	Sales	No. Observations:	36			
Model:	ARIMA(1, 0, 0)	Log Likelihood	-376.058			
Date:	Thu, 21 Sep 2023	AIC	758.116			
Time:	00:18:37	BIC	762.867			
Sample:	01-01-2014	HQIC	759.774			
	- 12-01-2016					
Covariance Type:	opg					
	coef	std err	z	P> z 	[0.025	0.975]
const	1.463e+04	1934.689	7.561	0.000	1.08e+04	1.84e+04
ar.L1	0.2825	0.170	1.658	0.097	-0.051	0.616
sigma2	6.929e+07	0.134	5.15e+08	0.000	6.93e+07	6.93e+07
Ljung-Box (L1) (Q):	0.01	Jarque-Bera (JB):	4.47			
Prob(Q):	0.92	Prob(JB):	0.11			
Heteroskedasticity (H):	1.40	Skew:	0.83			
Prob(H) (two-sided):	0.57	Kurtosis:	2.53			

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 5.41e+25. Standard errors may be unstable.

```
In [92]: start=len(train)
end=len(train)+len(test)-1
```

```
predictions = results.predict(start=start, end=end, dynamic=False, typ='levels').rename('ARIMA(1,0,0) Predictions')
predictions
```

```
Out[92]: 2017-01-01    20856.524035
         2017-02-01    16387.395968
         2017-03-01    15125.048311
         2017-04-01    14768.486230
         2017-05-01    14667.771885
         2017-06-01    14639.324162
         2017-07-01    14631.288832
         2017-08-01    14629.019177
         2017-09-01    14628.378091
         2017-10-01    14628.197011
         2017-11-01    14628.145863
         2017-12-01    14628.131415
         Freq: MS, Name: ARIMA(1,0,0) Predictions, dtype: float64
```


```
In [93]: for i in range(len(predictions)):
         print(f"predicted={predictions[i]:<11.10}, expected={test['Sales'][i]}")
```

```
predicted=20856.52404, expected=5964.032
predicted=16387.39597, expected=6866.3374
predicted=15125.04831, expected=10893.4448
predicted=14768.48623, expected=9065.9581
predicted=14667.77188, expected=16957.5582
predicted=14639.32416, expected=19008.5867
predicted=14631.28883, expected=11813.021999999999
predicted=14629.01918, expected=15441.874
predicted=14628.37809, expected=29028.206000000002
predicted=14628.19701, expected=21884.0682
predicted=14628.14586, expected=37056.715
predicted=14628.13142, expected=31407.4668
```

```
In [94]: title = 'Monthly Sales of Furnitures'
         ylabel='Sales Data'
         xlabel='Order Date'

         ax = test['Sales'].plot(legend=True,figsize=(12,6),title=title)
         predictions.plot(legend=True)
         ax.autoscale(axis='x',tight=True)
         ax.set(xlabel=xlabel, ylabel=ylabel)
```

```
ax.yaxis.set_major_formatter(formatter)
```



```
In [95]: from sklearn.metrics import mean_squared_error

error = mean_squared_error(test['Sales'], predictions)
print(f'ARIMA(1,0,0) MSE Error: {error:11.10}')
```

ARIMA(1,0,0) MSE Error: 120032385.7

```
In [96]: from statsmodels.tools.eval_measures import rmse
```




```
error = rmse(test['Sales'], predictions)
print(f'ARIMA(1,0,0) RMSE Error: {error:11.10}')
```

ARIMA(1,0,0) RMSE Error: 10955.92925

```
In [97]: model = ARIMA(df['Sales'],order=(1,1,1))
         results = model.fit()
         fcast = results.predict(len(df),len(df)+11,type='levels').rename('ARIMA(1,0,0) Forecast')
         title = 'Monthly Sales Data of Furniture'
         ylabel='Sales Data'
         xlabel='Order Date'

         ax = df['Sales'].plot(legend=True,figsize=(12,6),title=title)
         fcast.plot(legend=True)
         ax.autoscale(axis='x',tight=True)
         ax.set(xlabel=xlabel, ylabel=ylabel)
         ax.yaxis.set_major_formatter(formatter)
```



```
In [98]: import pandas as pd
import numpy as np
%matplotlib inline

from statsmodels.tsa.statespace.sarimax import SARIMAX

from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.seasonal import seasonal_decompose
from pmdarima import auto_arima

import warnings
```

```
warnings.filterwarnings("ignore")
```


```
In [99]: df = pd.read_excel('C://Users//saswa//OneDrive//Desktop//Pinaki-Time-series-forecasting//Superstore_Sales_Records.xls', index_
df = df[df['Category']=='Furniture']
df = df.groupby(by='Order Date').agg({'Sales':sum})
df.sort_index(inplace=True)
df.head(15)
```

Out[99]:

	Sales
Order Date	
2014-01-06	2573.820
2014-01-07	76.728
2014-01-10	51.940
2014-01-11	9.940
2014-01-13	879.939
2014-01-14	61.960
2014-01-16	127.104
2014-01-19	181.470
2014-01-20	1413.510
2014-01-21	25.248
2014-01-26	217.200
2014-01-27	333.000
2014-01-31	290.666
2014-02-08	14.560
2014-02-11	1650.050


```
In [100... df = df.resample('MS').sum()  
df.plot()
```


```
Out[100... <Axes: xlabel='Order Date'>
```



```
In [101... result = seasonal_decompose(df['Sales'], model='add')  
result.plot()
```

Out[101...





```
In [102... auto_arima(df['Sales'],seasonal=True,m=12).summary()
```

Out[102...

SARIMAX Results

Dep. Variable:	y			No. Observations:	48	
Model:	SARIMAX(2, 1, 0, 12)			Log Likelihood	-349.872	
Date:	Thu, 21 Sep 2023			AIC	707.744	
Time:	00:21:27			BIC	714.078	
Sample:	01-01-2014			HQIC	709.955	
	- 12-01-2017					
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975
intercept	1616.5115	1257.009	1.286	0.198	-847.181	4080.205
ar.S.L12	-0.1196	0.087	-1.378	0.168	-0.290	0.051
ar.S.L24	0.1214	0.098	1.240	0.215	-0.070	0.313
sigma2	1.862e+07	0.036	5.1e+08	0.000	1.86e+07	1.86e+07
Ljung-Box (L1) (Q):	0.71	Jarque-Bera (JB):	1.87			
Prob(Q):	0.40	Prob(JB):	0.39			
Heteroskedasticity (H):	0.96	Skew:	0.55			
Prob(H) (two-sided):	0.95	Kurtosis:	2.84			

Warnings:

- [1] Covariance matrix calculated using the outer product of gradients (complex-step).
- [2] Covariance matrix is singular or near-singular, with condition number 4.96e+24. Standard errors may be unstable.

```
In [103... train = df.iloc[:len(df)-6]
test = df.iloc[len(df)-6:]
model = SARIMAX(train['Sales'], order=(0,1,3), seasonal_order=(2,1,0,12))
results = model.fit()
results.summary()
```


Out[103...

SARIMAX Results

Dep. Variable:			Sales	No. Observations:			42
Model:			SARIMAX(0, 1, 3)x(2, 1, [], 12)		Log Likelihood		-289.730
Date:			Thu, 21 Sep 2023		AIC		591.461
Time:			00:21:48		BIC		599.664
Sample:			01-01-2014		HQIC		594.030
			- 06-01-2017				
Covariance Type:			opg				
	coef	std err	z	P> z	[0.025	0.975]	
ma.L1	-0.3976	0.146	-2.722	0.006	-0.684	-0.111	
ma.L2	-0.1639	0.156	-1.053	0.292	-0.469	0.141	
ma.L3	-0.1625	0.167	-0.972	0.331	-0.490	0.165	
ar.S.L12	-0.1664	0.118	-1.404	0.160	-0.399	0.066	
ar.S.L24	0.2224	0.130	1.710	0.087	-0.033	0.477	
sigma2	2.461e+07	5.24e-10	4.7e+16	0.000	2.46e+07	2.46e+07	
Ljung-Box (L1) (Q):		2.15	Jarque-Bera (JB):		0.87		
Prob(Q):		0.14	Prob(JB):		0.65		
Heteroskedasticity (H):		0.75	Skew:		0.17		
Prob(H) (two-sided):		0.65	Kurtosis:		2.22		

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 3.91e+32. Standard errors may be unstable.

```
In [104... start=len(train)
end=len(train)+len(test)-1
predictions = results.predict(start=start, end=end, dynamic=False, typ='levels').rename('SARIMA(0,1,3)(1,0,1,12) Predictions')
```


```
In [105... for i in range(len(predictions)):
    print(f"predicted={predictions[i]:<11.10}, expected={test['Sales'][i]}")
```

```
predicted=18982.66877, expected=11813.021999999999
predicted=15746.93545, expected=15441.874
predicted=29761.95657, expected=29028.206000000002
predicted=13953.797 , expected=21884.0682
predicted=35822.90086, expected=37056.715
predicted=34853.24127, expected=31407.4668
```

```
In [106... title = 'Monthly Sales of Furniture'
ylabel='Sales Data'
xlabel='Order Date'

ax = test['Sales'].plot(legend=True,figsize=(12,6),title=title)
predictions.plot(legend=True)
ax.autoscale(axis='x',tight=True)
ax.set(xlabel=xlabel, ylabel=ylabel)
```

```
Out[106... [Text(0.5, 0, 'Order Date'), Text(0, 0.5, 'Sales Data')]
```



```
In [107... from sklearn.metrics import mean_squared_error
error = mean_squared_error(test['Sales'], predictions)
print(f'SARIMA(2,0,1,12) MSE Error: {error:11.10}')
```


SARIMA(2,0,1,12) MSE Error: 21386691.26

```
In [108... model = SARIMAX(df['Sales'],order=(1,0,2),seasonal_order=(2,1,0,12))
results = model.fit()
fcast = results.predict(len(df),len(df)+12,type='levels').rename('SARIMA(0,1,3)(1,0,1,12) Forecast')
```

```
In [109... title = 'Monthly Sales of Furniture'
ylabel='Sales Data'
xlabel='Order Date'

ax = df['Sales'].plot(legend=True,figsize=(12,6),title=title)
fcast.plot(legend=True)
ax.autoscale(axis='x',tight=True)
ax.set(xlabel=xlabel, ylabel=ylabel)
```

```
Out[109... [Text(0.5, 0, 'Order Date'), Text(0, 0.5, 'Sales Data')]
```



```
In [110... import pandas as pd
import numpy as np
%matplotlib inline

from statsmodels.tsa.statespace.sarimax import SARIMAX

from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.seasonal import seasonal_decompose
from pmdarima import auto_arima
import warnings
warnings.filterwarnings("ignore")

In [111... df = pd.read_excel('C://Users//saswa//OneDrive//Desktop//Pinaki-Time-series-forecasting//Superstore_Sales_Records.xls', index_
df = df[df['Category']=='Furniture']
df = df.groupby(by='Order Date').agg({'Sales':sum, 'Quantity':sum})
df.sort_index(inplace=True)
df.head(24)
```

Out[111...

	Sales	Quantity
Order Date		
2014-01-06	2573.820	9
2014-01-07	76.728	3
2014-01-10	51.940	1
2014-01-11	9.940	2
2014-01-13	879.939	9
2014-01-14	61.960	4
2014-01-16	127.104	6
2014-01-19	181.470	5
2014-01-20	1413.510	15
2014-01-21	25.248	3
2014-01-26	217.200	8
2014-01-27	333.000	3
2014-01-31	290.666	2
2014-02-08	14.560	2
2014-02-11	1650.050	10
2014-02-12	129.568	2
2014-02-18	25.160	5
2014-02-20	20.320	4
2014-03-01	1893.995	23
2014-03-03	928.802	8
2014-03-07	966.984	9

	Sales	Quantity
Order Date		
2014-03-11	8.320	5
2014-03-14	1139.920	4
2014-03-15	45.696	3


```
In [112... df = df.resample('W').sum()

df.head()
```

```
Out[112...      Sales  Quantity
Order Date
2014-01-12  2712.428      15
2014-01-19  1250.473      24
2014-01-26  1655.958      26
2014-02-02   623.666       5
2014-02-09   14.560       2
```


```
In [113... df['Sales'].plot()
```

```
Out[113... <Axes: xlabel='Order Date'>
```




```
In [114... df['Quantity'].plot()
```


```
Out[114... <Axes: xlabel='Order Date'>
```

```
In [115... result_sales = seasonal_decompose(df['Sales'])  
result_sales.plot()
```


Out[115...






```
In [116... result_sales = seasonal_decompose(df['Quantity'])  
result_sales.plot()
```

Out[116...





In [117...

```

from statsmodels.tsa.stattools import adfuller

def adf_test(series,title=''):
    """
    Pass in a time series and an optional title, returns an ADF report
    """
    print(f'Augmented Dickey-Fuller Test: {title}')
    result = adfuller(series.dropna(),autolag='AIC') # .dropna() handles differenced data

    labels = ['ADF test statistic','p-value','# lags used','# observations']
    out = pd.Series(result[0:4],index=labels)

```

```

for key,val in result[4].items():
    out[f'critical value ({key})']=val

print(out.to_string())          # .to_string() removes the line "dtype: float64"

if result[1] <= 0.05:
    print("Strong evidence against the null hypothesis")
    print("Reject the null hypothesis")
    print("Data has no unit root and is stationary")
else:
    print("Weak evidence against the null hypothesis")
    print("Fail to reject the null hypothesis")
    print("Data has a unit root and is non-stationary")
adf_test(df['Sales'])

```

Augmented Dickey-Fuller Test:

```

ADF test statistic      -3.444275
p-value                0.009539
# lags used            6.000000
# observations         201.000000
critical value (1%)    -3.463309
critical value (5%)    -2.876029
critical value (10%)   -2.574493
Strong evidence against the null hypothesis
Reject the null hypothesis
Data has no unit root and is stationary

```

In [118...

```
adf_test(df['Quantity'])
```

Augmented Dickey-Fuller Test:

```

ADF test statistic      -4.457632
p-value                0.000234
# lags used            4.000000
# observations         203.000000
critical value (1%)    -3.462980
critical value (5%)    -2.875885
critical value (10%)   -2.574416
Strong evidence against the null hypothesis
Reject the null hypothesis
Data has no unit root and is stationary

```

```
In [119... auto_arima(df['Sales'], bseasonal=True, m=12).summary()
```

```
Out[119...
```

SARIMAX Results

Dep. Variable:	y	No. Observations:	208			
Model:	SARIMAX(1, 1, 2)	Log Likelihood	-1917.198			
Date:	Thu, 21 Sep 2023	AIC	3842.396			
Time:	00:26:55	BIC	3855.727			
Sample:	01-12-2014	HQIC	3847.787			
	- 12-31-2017					
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.7864	0.166	4.731	0.000	0.461	1.112
ma.L1	-1.4945	0.201	-7.446	0.000	-1.888	-1.101
ma.L2	0.5044	0.183	2.763	0.006	0.147	0.862
sigma2	6.442e+06	4.47e-08	1.44e+14	0.000	6.44e+06	6.44e+06
Ljung-Box (L1) (Q):	0.12	Jarque-Bera (JB):	130.12			
Prob(Q):	0.73	Prob(JB):	0.00			
Heteroskedasticity (H):	1.43	Skew:	1.32			
Prob(H) (two-sided):	0.14	Kurtosis:	5.85			

Warnings:

- [1] Covariance matrix calculated using the outer product of gradients (complex-step).
- [2] Covariance matrix is singular or near-singular, with condition number 7.8e+29. Standard errors may be unstable.

In [120... `len(df)`

Out[120... 208

In [121... `train = df.iloc[:170]`
`test = df.iloc[170:]`

In [122... `model = SARIMAX(train['Sales'],order=(1,1,2) ,enforce_invertibility=False)`
`results = model.fit()`
`results.summary()`

Out [122...

SARIMAX Results

Dep. Variable:	Sales	No. Observations:	170
Model:	SARIMAX(1, 1, 2)	Log Likelihood	-1573.035
Date:	Thu, 21 Sep 2023	AIC	3154.069
Time:	00:27:21	BIC	3166.589
Sample:	01-12-2014	HQIC	3159.150
	- 04-09-2017		
Covariance Type:	opg		
	coef	std err	z P> z [0.025 0.975]
ar.L1	-0.7910	0.528	-1.499 0.134 -1.825 0.243
ma.L1	-0.2794	0.479	-0.583 0.560 -1.218 0.660
ma.L2	-0.9240	0.525	-1.759 0.079 -1.953 0.105
sigma2	6.972e+06	9.33e+05	7.472 0.000 5.14e+06 8.8e+06
Ljung-Box (L1) (Q):	5.41	Jarque-Bera (JB):	70.94
Prob(Q):	0.02	Prob(JB):	0.00
Heteroskedasticity (H):	1.23	Skew:	1.14
Prob(H) (two-sided):	0.44	Kurtosis:	5.21

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).


In [123...

```
start=len(train)
end=len(train)+len(test)-1
```

```
predictions = results.predict(start=start, end=end, dynamic=False).rename('SARIMA(1,0,0)(2,0,0,7) Predictions')
title='Restaurant Visitors'
ylabel='Visitors per day'
xlabel=''

ax = test['Sales'].plot(legend=True,figsize=(12,6),title=title)
predictions.plot(legend=True)
ax.autoscale(axis='x',tight=True)
ax.set(xlabel=xlabel, ylabel=ylabel)
```

Out[123... [Text(0.5, 0, ''), Text(0, 0.5, 'Visitors per day')]



```
In [124... from statsmodels.tools.eval_measures import mse,rmse

error1 = mse(test['Sales'], predictions)
error2 = rmse(test['Sales'], predictions)

print(f'SARIMA(1,0,0)(2,0,0,7) MSE Error: {error1:11.10}')
print(f'SARIMA(1,0,0)(2,0,0,7) RMSE Error: {error2:11.10}')
```

SARIMA(1,0,0)(2,0,0,7) MSE Error: 12077452.14
SARIMA(1,0,0)(2,0,0,7) RMSE Error: 3475.262888

```
In [125... df.columns
```

```
Out[125... Index(['Sales', 'Quantity'], dtype='object')
```

```
In [126... model = SARIMAX(train['Sales'],exog=train['Quantity'],order=(1,0,0),seasonal_order=(2,0,0,7),enforce_invertibility=False)
results = model.fit()
results.summary()
```

Out[126...

SARIMAX Results

Dep. Variable:		Sales		No. Observations:		170
Model:		SARIMAX(1, 0, 0)x(2, 0, 0, 7)		Log Likelihood		-1466.114
Date:		Thu, 21 Sep 2023		AIC		2942.229
Time:		00:28:10		BIC		2957.908
Sample:		01-12-2014		HQIC		2948.591
		- 04-09-2017				
Covariance Type:		opg				
	coef	std err	z	P> z	[0.025	0.975]
Quantity	97.6600	2.403	40.643	0.000	92.950	102.369
ar.L1	-0.0783	0.072	-1.086	0.277	-0.220	0.063
ar.S.L7	0.2222	0.057	3.874	0.000	0.110	0.335
ar.S.L14	-0.0828	0.094	-0.879	0.379	-0.267	0.102
sigma2	1.903e+06	1.58e+05	12.024	0.000	1.59e+06	2.21e+06
Ljung-Box (L1) (Q):		0.02	Jarque-Bera (JB):		295.39	
Prob(Q):		0.88	Prob(JB):		0.00	
Heteroskedasticity (H):		1.16	Skew:		1.56	
Prob(H) (two-sided):		0.58	Kurtosis:		8.66	


Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [127... start=len(train)
end=len(train)+len(test)-1
exog_forecast = test[['Quantity']]
predictions = results.predict(start=start, end=end, exog=exog_forecast).rename('SARIMAX(1,0,0)(2,0,0,7) Predictions')
title='Restaurant Visitors'
ylabel='Visitors per day'
xlabel=''

ax = test['Sales'].plot(legend=True,figsize=(12,6),title=title)
predictions.plot(legend=True)
ax.autoscale(axis='x',tight=True)
ax.set(xlabel=xlabel, ylabel=ylabel)
```

```
Out[127... [Text(0.5, 0, ''), Text(0, 0.5, 'Visitors per day')]
```



In [128...

```
print(f'SARIMA(1,0,0)(2,0,0,7) MSE Error: {error1:11.10}')
```

```
print(f'SARIMA(1,0,0)(2,0,0,7) RMSE Error: {error2:11.10}')
```

```
print()
```

```
error1x = mse(test['Sales'], predictions)
```

```
error2x = rmse(test['Sales'], predictions)
```



```
print(f'SARIMAX(1,0,0)(2,0,0,7) MSE Error: {error1x:11.10}')
```

```
print(f'SARIMAX(1,0,0)(2,0,0,7) RMSE Error: {error2x:11.10}')
```

```
SARIMA(1,0,0)(2,0,0,7) MSE Error: 12077452.14  
SARIMA(1,0,0)(2,0,0,7) RMSE Error: 3475.262888
```


```
SARIMAX(1,0,0)(2,0,0,7) MSE Error: 2317778.515  
SARIMAX(1,0,0)(2,0,0,7) RMSE Error: 1522.425208
```

```
In [129... model = SARIMAX(df['Sales'],exog=df['Quantity'],order=(1,2,1),seasonal_order=(1,2,2,7),enforce_invertibility=False)  
           results = model.fit()  
           exog_forecast = df[169:][['Quantity']]  
           fcast = results.predict(len(df),len(df)+38,exog=exog_forecast).rename('SARIMAX(1,0,0)(2,0,0,7) Forecast')
```

```
In [130... title='Restaurant Visitors'  
          ylabel='Visitors per day'  
          xlabel=''
```

```
ax = df['Sales'].plot(legend=True,figsize=(16,6),title=title)  
fcast.plot(legend=True)  
ax.autoscale(axis='x',tight=True)  
ax.set(xlabel=xlabel, ylabel=ylabel)
```

```
Out[130... [Text(0.5, 0, ''), Text(0, 0.5, 'Visitors per day')]
```



In []: